



Model Evaluation

20 Oct 2566

Sarun Gulyanon

Agenda

- Model Evaluation
- Case Study : Kaggle

The top corners of the slide feature decorative circuit-like patterns. These consist of thin, light blue lines that branch out and connect to small, solid blue circular nodes, resembling a stylized electronic circuit board.

1.

Model Evaluation

Recap: Generalization

- The error rate on new cases is called the *generalization error* (or *out-of-sample error*)
- This value tells you how well your model will perform on instances it has **never seen before**.
- We can only estimate of this error.



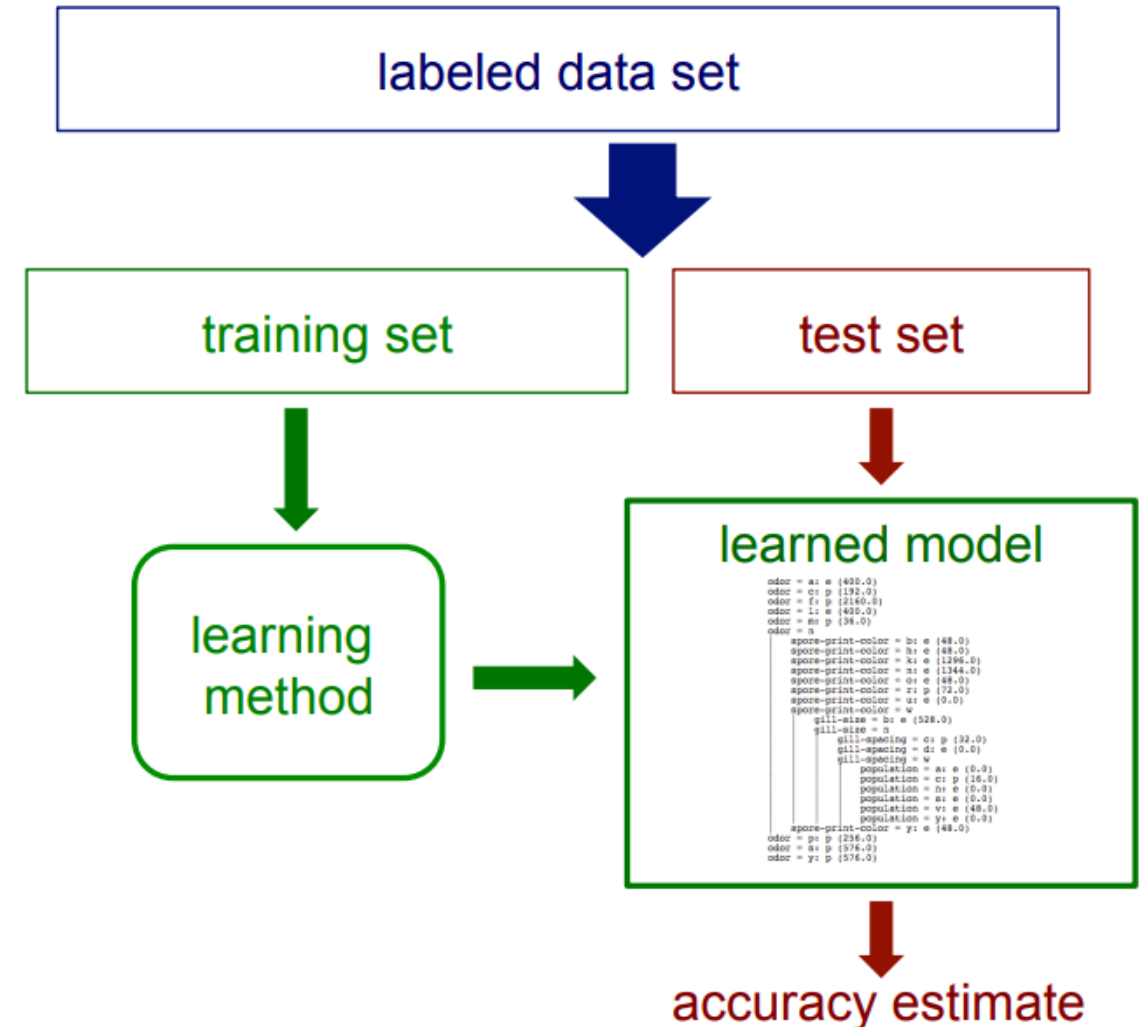
Model Evaluation

- Methods for Performance Evaluation
 - How to obtain reliable estimates?
- Metrics for Performance Evaluation
 - How to evaluate the performance of a model?
- Methods for Model Comparison
 - How to compare the relative performance among competing models?

Test Set

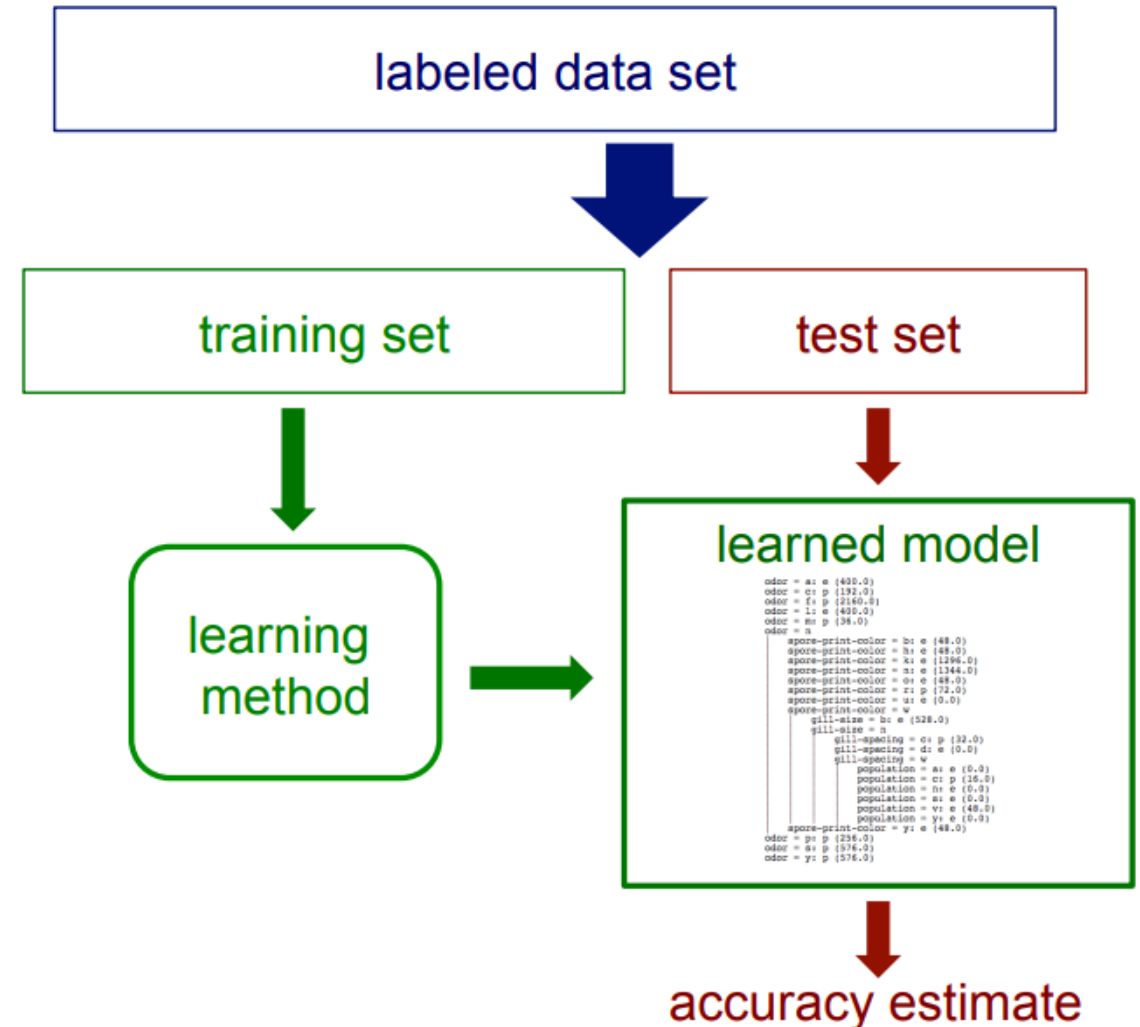
Hold-Off Evaluation

- How can we get an unbiased estimate of the accuracy of a learned model?
- Split data into two sets: **training** and **test** sets.
- Build model using the training set
- Evaluating model on the test set gives an estimation of generalization error

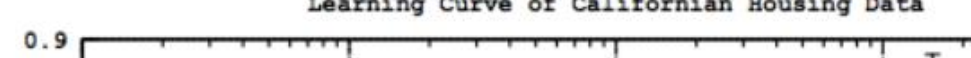



Test Set

- How can we get an unbiased estimate of the accuracy of a learned model?
- If the test-set labels influence the learned model in any way, accuracy estimates will be biased




Learning Curves

- How does the accuracy of a learning method change as a function of the training-set size?
 - Plot accuracy against the training-set size.
- 
- The graph shows the learning curve for the Californian Housing Data. The y-axis represents accuracy, ranging from 0.86 to 0.9. The x-axis represents the training-set size, with major ticks at 10,000, 20,000, 30,000, and 40,000. A solid line with error bars shows the accuracy increasing from approximately 0.855 at 10,000 to 0.89 at 40,000. The error bars indicate the variability in accuracy across different runs.
- | Training-set size | Accuracy (approx.) |
|-------------------|--------------------|
| 10,000 | 0.855 |
| 20,000 | 0.865 |
| 30,000 | 0.878 |
| 40,000 | 0.890 |



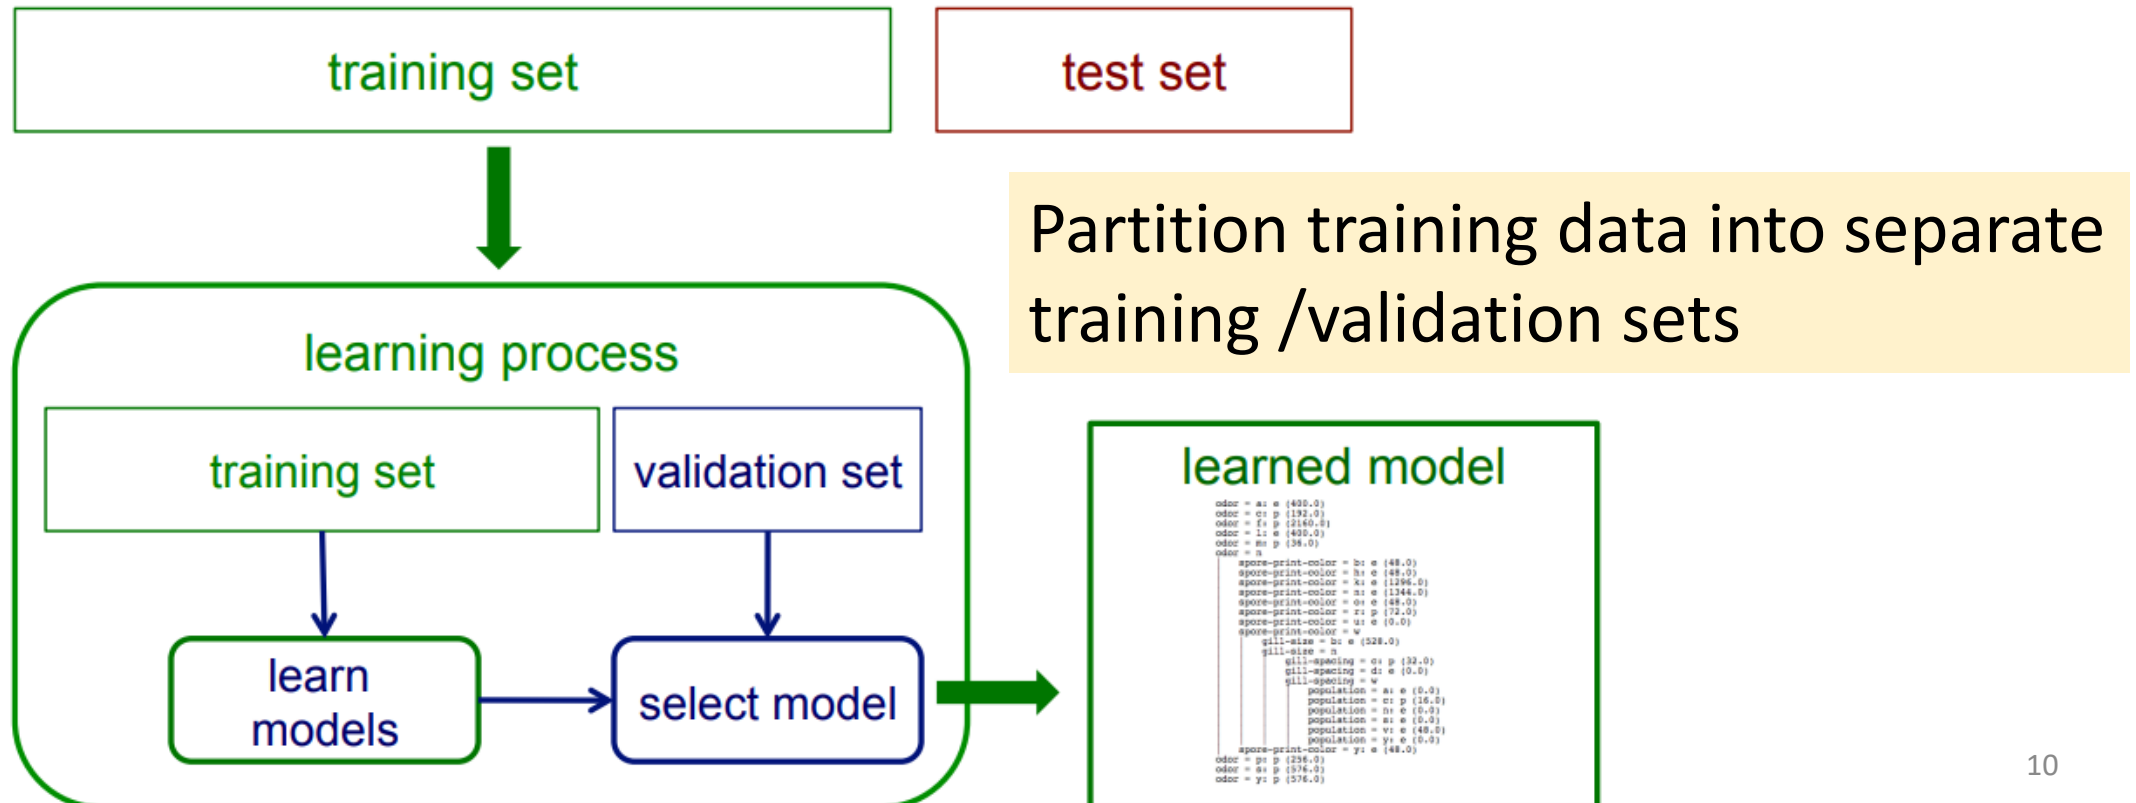
Learning Curves

- Given training/test set partition
- For each sample size s , randomly select s instances from training set
- Learn model
- Evaluate model on test set to determine accuracy a
- plot (s, a) or repeat n times and plot $(s, \text{avg. accuracy and error bars})$




Validation Set

- Suppose we want unbiased estimates of accuracy during the learning process (e.g., to choose the best level of decision-tree pruning)?



Fine-Tune Your Model

- Find the best hyperparameter for the model
 1. Grid Search = Search through all combinations of the parameters. (Scikit-Learn's GridSearch)
 2. Randomized Search = evaluates a preset number of random combinations. (Scikit-Learn's RandomizedSearchCV)



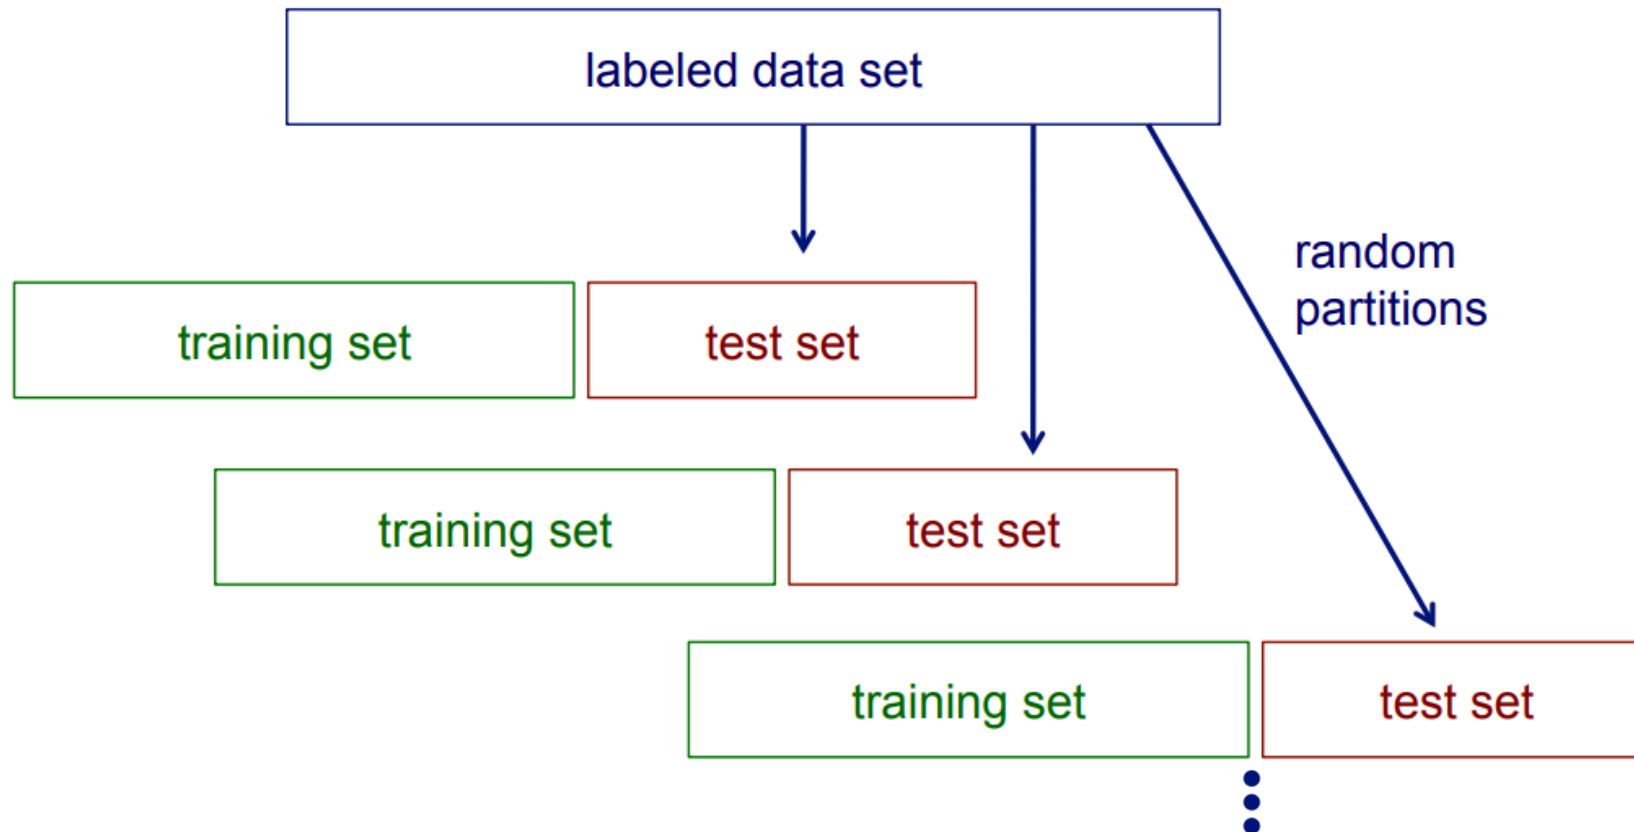
<https://medium.com/@cjl2fv/an-intro-to-hyper-parameter-optimization-using-grid-search-and-random-search-d73b9834ca0a>

Limitations of Using a Single Training/Test Partition

- We may not have enough data to make sufficiently large training and test sets
 - A larger test set gives us more reliable estimate of accuracy (i.e. a lower variance estimate)
 - But... a larger training set will be more representative of how much data we actually have for learning process
- A single training set doesn't tell us how sensitive accuracy is to a particular training sample

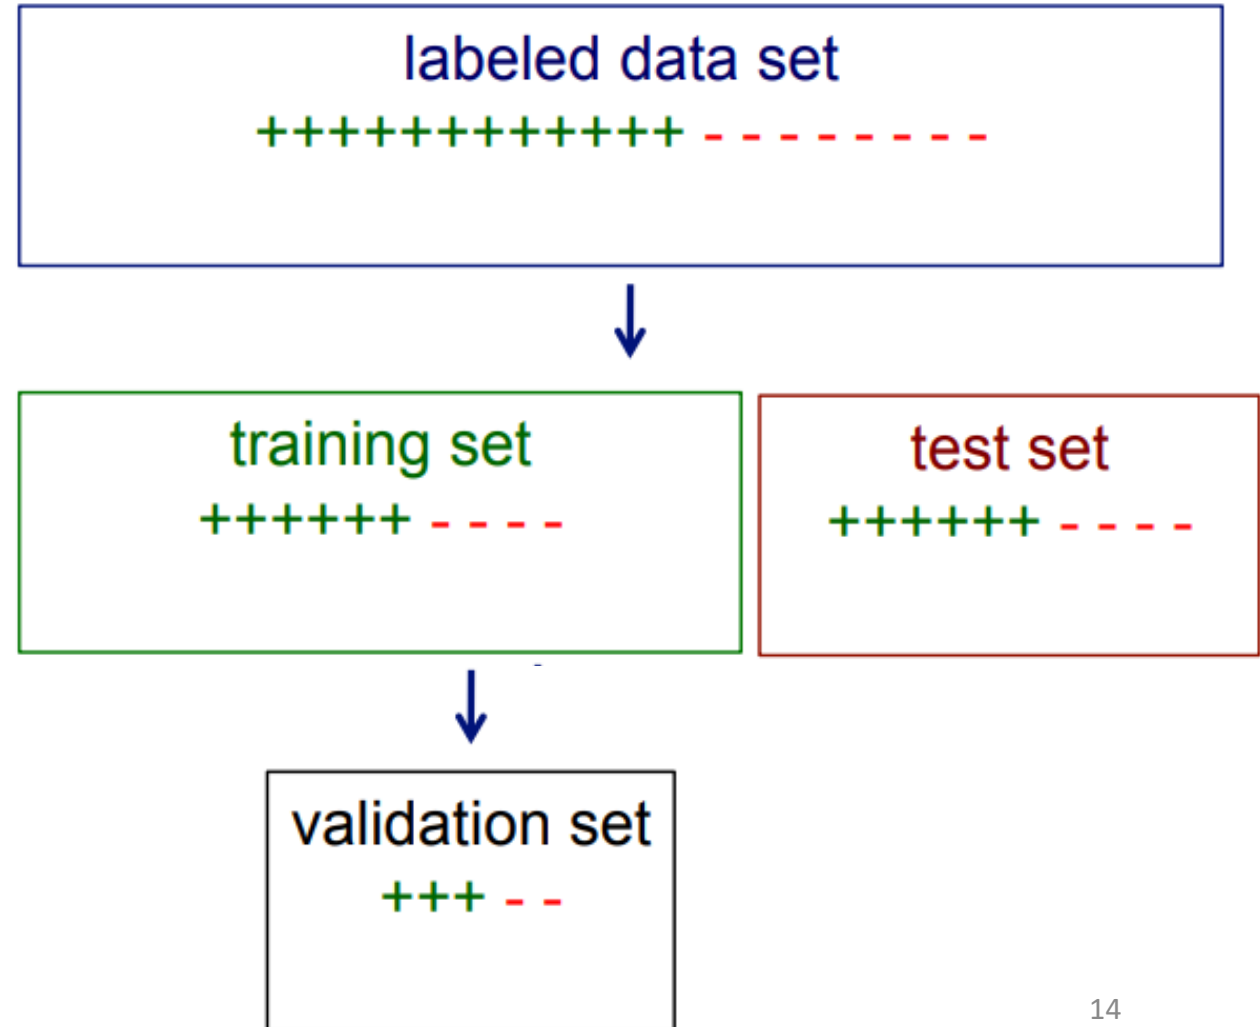
Random Resampling

- We can address the second issue by repeatedly randomly partitioning the available data into training and test sets.



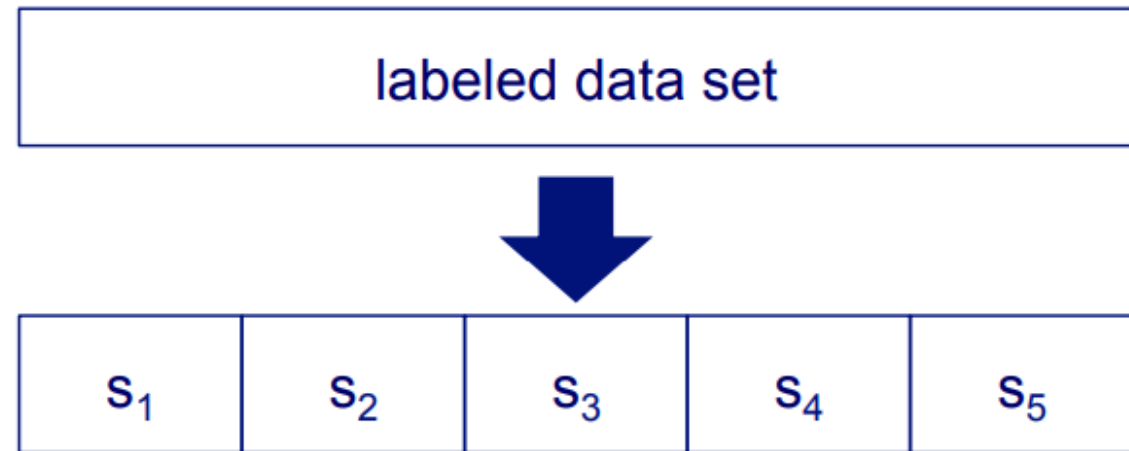
Stratified Sampling

- When randomly selecting training or validation sets, we may want to ensure that class proportions are maintained in each selected set.
- This can be done via stratified sampling: first stratify instances by class, then randomly select instances from each class proportionally



Cross Validation

- Partition data into n subsamples
- Iteratively leave one subsample out for the test set, train on the rest
- 10-fold CV is common,
- But smaller values of n are often used when learning takes time



| iteration | train on | test on |
|-----------|-------------------------|---------|
| 1 | s_2 s_3 s_4 s_5 | s_1 |
| 2 | s_1 s_3 s_4 s_5 | s_2 |
| 3 | s_1 s_2 s_4 s_5 | s_3 |
| 4 | s_1 s_2 s_3 s_5 | s_4 |
| 5 | s_1 s_2 s_3 s_4 | s_5 |

Cross Validation Example

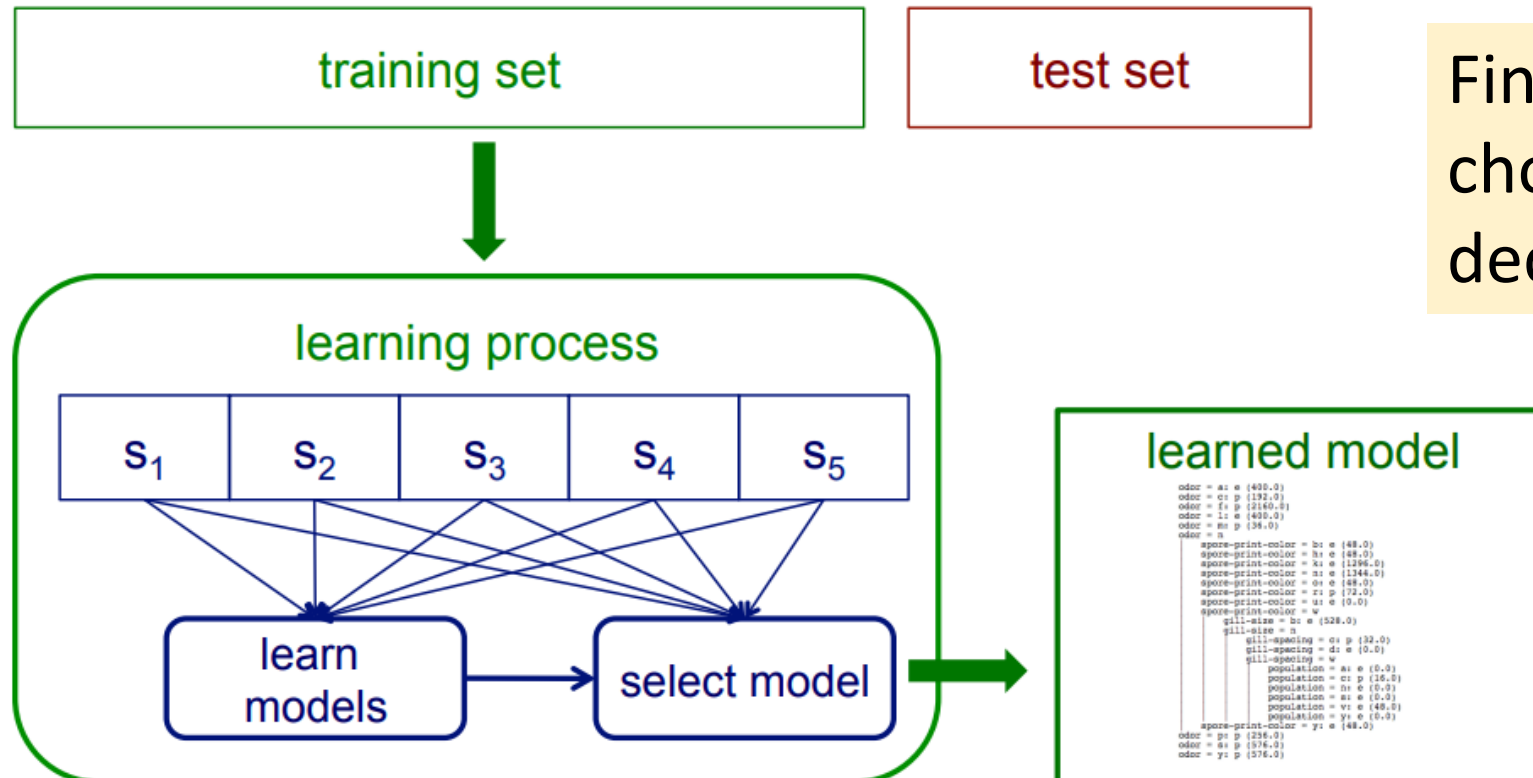
- Leave-one-out cross validation, $n = \#$ instances
- CV makes efficient use of the available data for testing
- Whenever we use multiple training sets, we are evaluating a learning method as opposed to an individual learned model.

| iteration | train on | test on | correct |
|-----------|-------------------------|---------|---------|
| 1 | s_2 s_3 s_4 s_5 | s_1 | 11 / 20 |
| 2 | s_1 s_3 s_4 s_5 | s_2 | 17 / 20 |
| 3 | s_1 s_2 s_4 s_5 | s_3 | 16 / 20 |
| 4 | s_1 s_2 s_3 s_5 | s_4 | 13 / 20 |
| 5 | s_1 s_2 s_3 s_4 | s_5 | 16 / 20 |

Accuracy = $73/100 = 73\%$

Internal Cross Validation

- Instead of a single validation set, we can use cross-validation within a training set to select a model



Fine tuning model, e.g., to choose the best level of decision-tree pruning

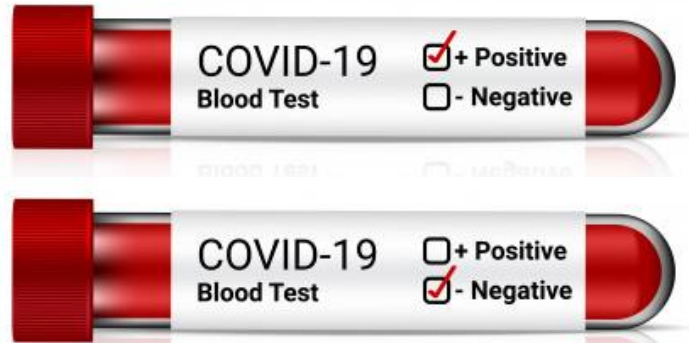
Performance Measures

- Need a metric to measure/evaluate/compare classifiers and regressors.
- Evaluating classifier is trickier than evaluating a regressor.
→ real work is continuous, discrete attributes are what we defined.
- Metrics
 - **Classification:** Accuracy, F1-score, precision/recall, ROC curve
 - **Regression:** MAE, RMSE, R^2

Terminology

| | | actual class | |
|-----------------|----------|----------------------|----------------------|
| | | positive | negative |
| predicted class | positive | true positives (TP) | false positives (FP) |
| | negative | false negatives (FN) | true negatives (TN) |

pos
neg



pred *GT*

| | | |
|-----------|----------|----------|
| <i>TP</i> | <i>+</i> | <i>+</i> |
| <i>FP</i> | <i>+</i> | <i>-</i> |
| <i>FN</i> | <i>-</i> | <i>+</i> |
| <i>TN</i> | <i>-</i> | <i>-</i> |

Accuracy

- The most simple metric for classification
- correctly predicted samples / total testing samples

$$\text{Accuracy} = \frac{TP + TN}{TP + FP + FN + TN}$$

- Example:

$$\text{Accuracy} = 6/8 = 75\%$$




| | Ground Truth | Prediction |
|---|--------------|------------|
| 1 | Apple | Apple |
| 2 | Papaya | Papaya |
| 3 | Papaya | Papaya |
| 4 | Orange | Orange |
| 5 | Apple | Apple |
| 6 | Papaya | Orange |
| 7 | Apple | Apple |
| 8 | Banana | Papaya |

Accuracy Issues

- Accuracy may not be useful measure in cases where:
 - ❑ There is a large class skew
 - Is 98% accuracy good if 97% of the instances are negative?
 - ❑ There are differential misclassification costs – say, getting a positive wrong costs more than getting a negative wrong
 - Consider a medical domain in which a false positive results in an extraneous test but a false negative results in a failure to treat a disease

Confusion Matrix

- Or error matrix
- A table whose entry (A,B) indicates the number of times instances of class A are predicted as class B.




Precision and Recall


$n = 100$

Predict class 1 60
Ground Truth Class 1 40

How many selected items are relevant?



How many relevant items are selected?



$$\frac{TP}{TP + FP}$$

Example

Neg Class 1 99
Pos 2 1

$P = 0$ $R = 0$
Automated Diagnosis

- Automated system that diagnoses patients either he/she is unwell/normal TP
- Here, Positives = Unwell

$$\text{Precision} = \frac{TP}{TP + FP} = \frac{1}{1+3} = \underline{0.25}$$

$$\text{Recall} = \frac{TP}{TP + FN} = \frac{1}{1+2} = \underline{0.33}$$

| | Ground Truth | Prediction |
|---|--------------|------------|
| 1 | Unwell ✓ | Unwell ✓ |
| 2 | Unwell ✓ | Normal |
| 3 | Unwell ✓ | Normal |
| 4 | Normal | Normal |
| 5 | Normal | Unwell ✓ |
| 6 | Normal | Unwell ✓ |
| 7 | Normal | Unwell ✓ |
| 8 | Normal | Normal |


F1-Score

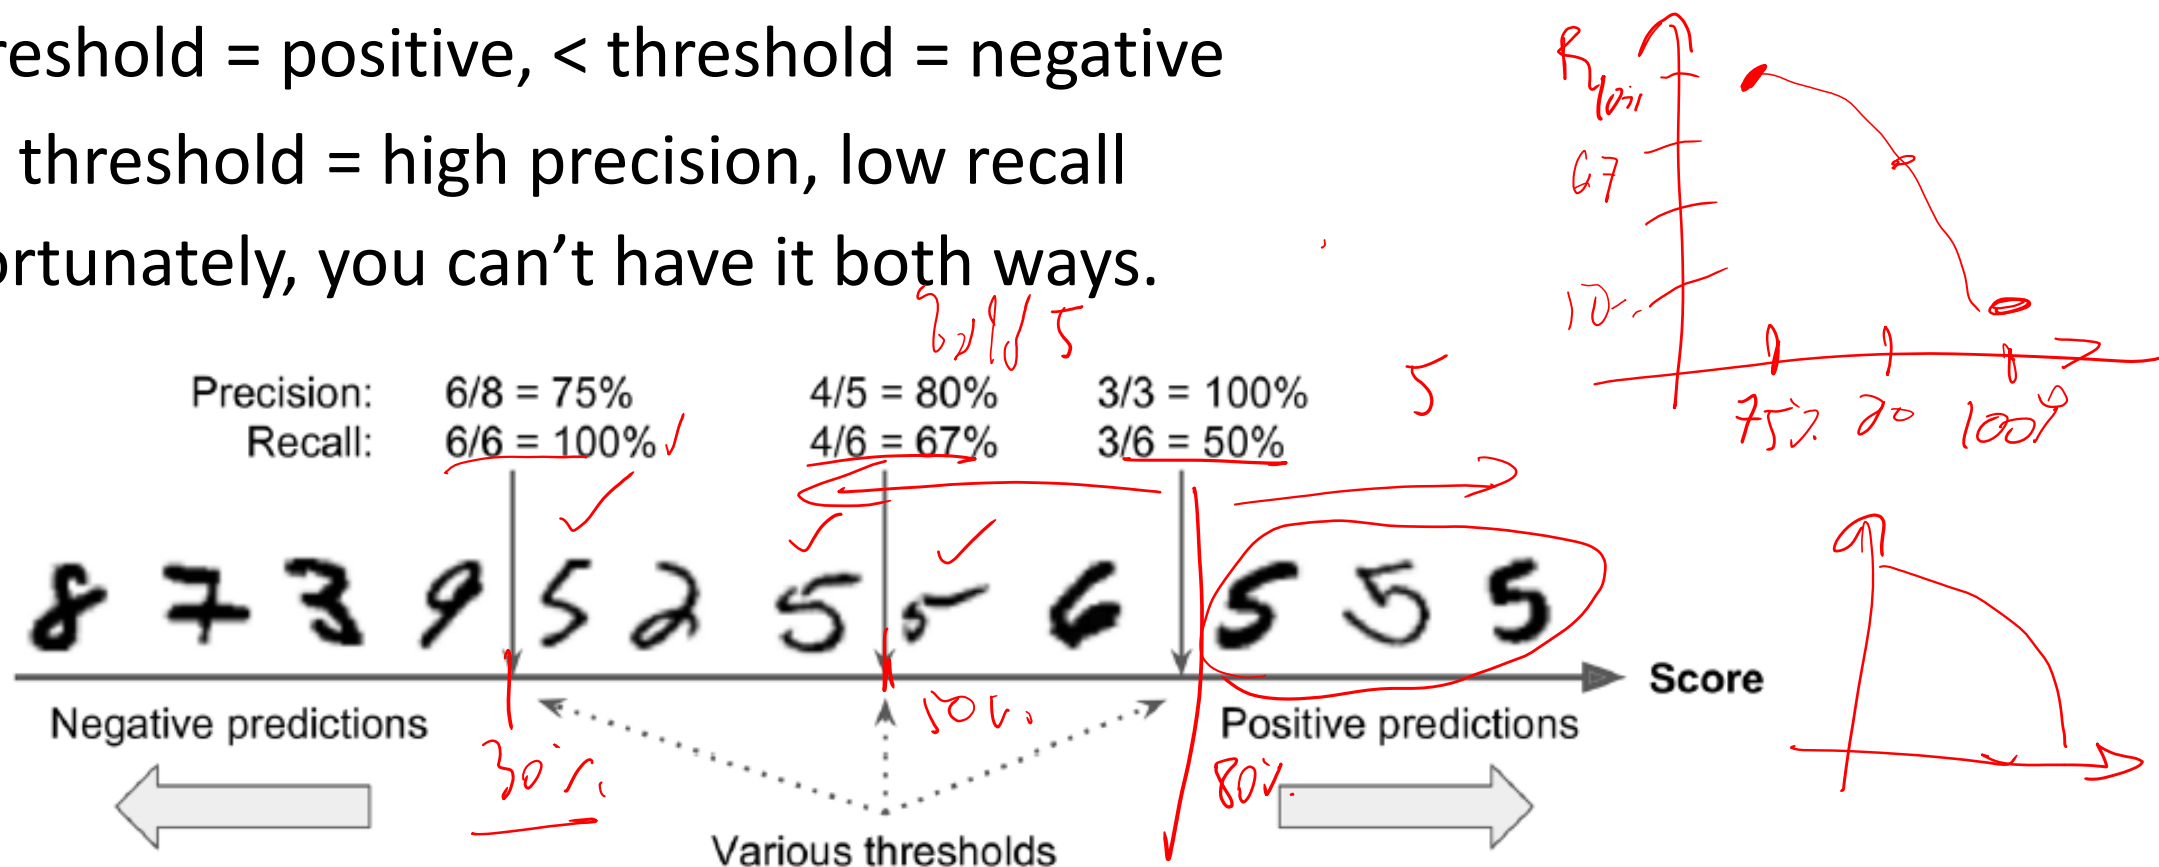
- F1 score = harmonic mean of precision and recall.

$$F_1 = \frac{2}{\frac{1}{\text{precision}} + \frac{1}{\text{recall}}} = 2 \times \frac{\text{precision} \times \text{recall}}{\text{precision} + \text{recall}}$$

- The harmonic mean gives much more weight to low values.
→ Give high F1 score if both recall and precision are high.

Precision/Recall Tradeoff

- Assume a classifier is fixed and it assigns a score to each sample
 - $> \text{threshold} = \text{positive}$, $< \text{threshold} = \text{negative}$
 - High threshold = high precision, low recall
 - Unfortunately, you can't have it both ways.
- 
- A hand-drawn red diagram on the right side of the slide. It consists of a vertical line with several horizontal tick marks. To the left of the top tick mark is the label
- $y_{(27)}$
- . To the left of the second tick mark from the top is the label
- $G7$
- . There are also some red dots and a small red arrow pointing upwards near the top of the line.



ROC Curve

- Receiver Operating Characteristic (ROC) curve
- For binary classifiers
- Plots the true positive rate (another name for recall) against the false positive rate for various threshold values.

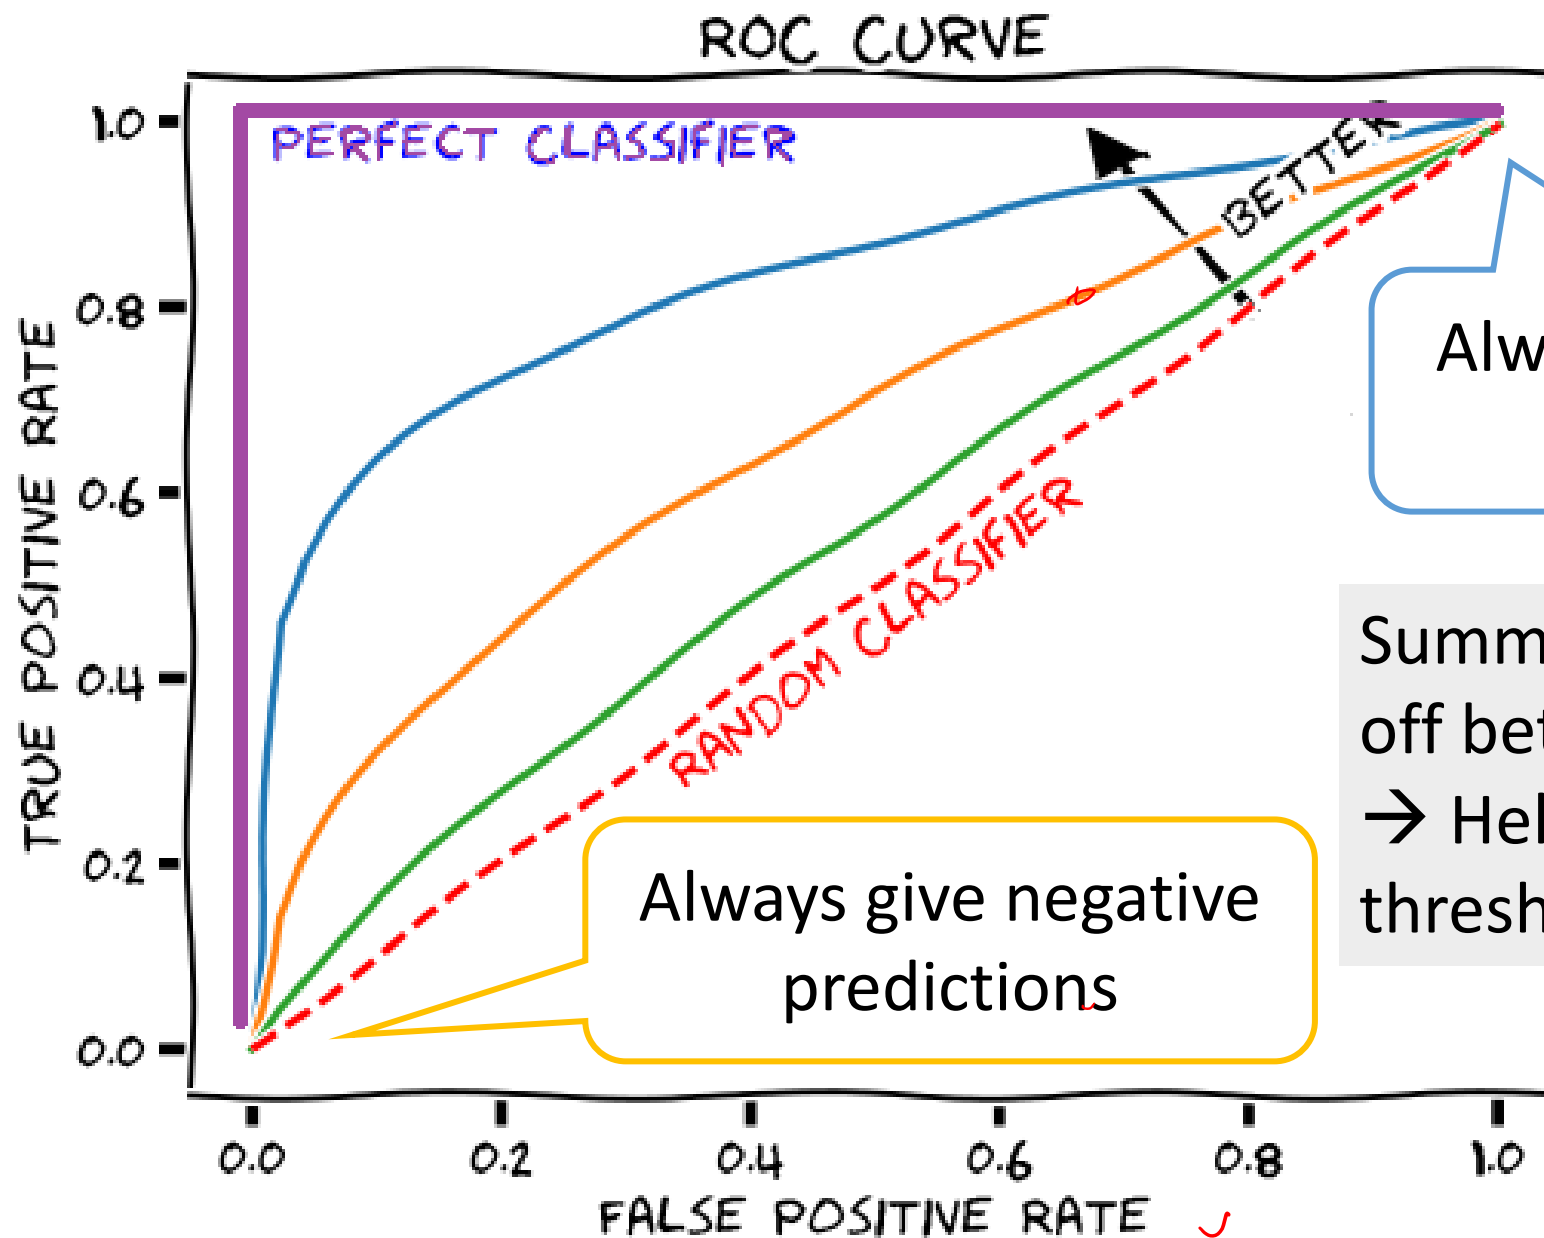
$$\text{TPR} = \frac{\text{TP}}{\text{P}} = \frac{\text{TP}}{\text{TP} + \text{FN}}$$

recall

$$\text{FPR} = \frac{\text{FP}}{\text{N}} = \frac{\text{FP}}{\text{FP} + \text{TN}}$$

- TPR = probability of detection
- FPR = probability of false alarm
- <http://arogozhnikov.github.io/RocCurve.html>

Recall




Always give positive predictions

Summarize the trade-off between TPR & FPR
→ Help determine the threshold

Always give negative predictions

ROC Curve Example

| | instance | confidence positive | correct class |
|----------|----------|---------------------|---------------|
| 100% | Ex 9 | .99 | + |
| 90% | Ex 7 | .98 | + |
| | | TPR= 2/5, FPR= 0/5 | |
| 70% | Ex 1 | .72 | - |
| | | TPR= 2/5, FPR= 1/5 | |
| | Ex 2 | .70 | + |
| 60% | Ex 6 | .65 | + |
| | | TPR= 4/5, FPR= 1/5 | |
| | Ex 10 | .51 | - |
| 30% | Ex 3 | .39 | - |
| | | TPR= 4/5, FPR= 3/5 | |
| | Ex 5 | .24 | + |
| | | TPR= 5/5, FPR= 3/5 | |
| 20% 100% | Ex 4 | .11 | - |
| 0% | Ex 8 | .01 | - |
| | | TPR= 5/5, FPR= 5/5 | |



AUC

- To compare classifiers, we need a number.
→ Measure the area under the curve (AUC) of ROC.
- A perfect classifier will have ROC AUC = 1
- A purely random classifier will have ROC AUC = 0.5.
- To compute, one possible solution similar to Trapezoidal rule, divide regions at thresholds.

What about multiclass problem?

Demo: https://colab.research.google.com/drive/1u9AEoqBkhZDtyvI82OgG4jskN_9qkF2L?usp=sharing

Multilabel Classification

- Classification task labelling each sample with x labels from $n_classes$ possible classes, where x can be 0 to $n_classes$ inclusive.
- Thus, comparable to running $n_classes$ binary classification tasks
- Multilabel classifiers may treat the multiple classes simultaneously, accounting for correlated behavior among them.
- E.g., In diagnosis system, one patient may have multiple diseases.

RMSE

- A typical performance measure for regression problems.
- Measure the standard deviation of the errors the system makes in its predictions

$$\text{RMSE}(\mathbf{X}, h) = \sqrt{\frac{1}{m} \sum_{i=1}^m \left(h(\mathbf{x}^{(i)}) - y^{(i)} \right)^2}$$

- y = the expected output,
- \mathbf{X} = data, m = no. of testing data and
- h = prediction function

MAE

- Or Average Absolute Deviation
- Another popular performance measure for regression problems

$$\text{MAE}(\mathbf{X}, h) = \frac{1}{m} \sum_{i=1}^m \left| h(\mathbf{x}^{(i)}) - y^{(i)} \right|$$

- Since errors are squared in RMSE so it gives a relatively high weight to large errors
→ MAE is better when outliers are expected.

R^2

- Or coefficient of determination
- How well observed outcomes are replicated by the model →
The proportion of total variation of outcomes explained by the model.

$$R^2 = 1 - \frac{SS_{\text{res}}}{SS_{\text{tot}}}$$

Total sum of squares

$$SS_{\text{tot}} = \sum_i (y_i - \bar{y})^2$$

Observed data

Mean of the observed data

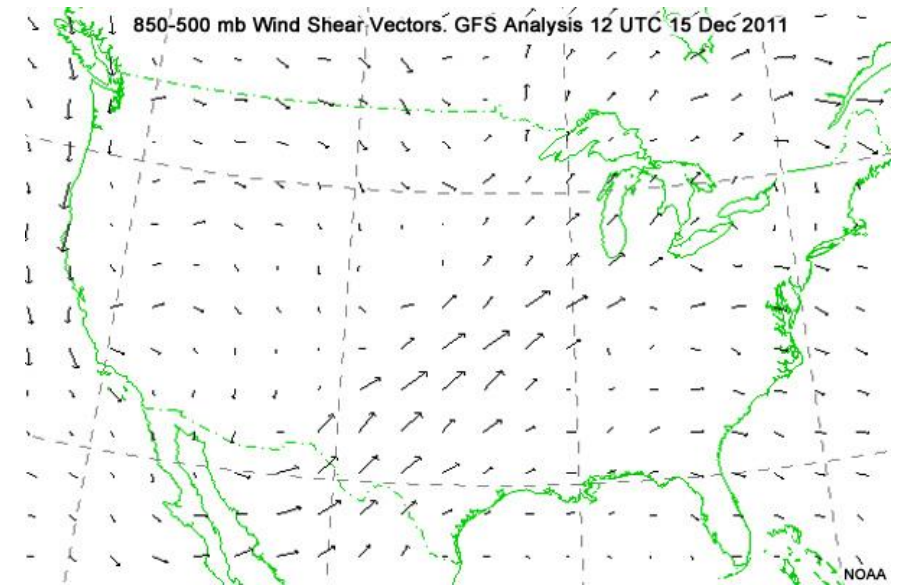
Residual sum of squares

$$SS_{\text{res}} = \sum_i (y_i - f_i)^2 = \sum_i e_i^2$$

Predicted value

Multioutput Regression

- Predicts multiple numerical properties for each sample.
- Thus comparable to running `n_output` regression estimators
- Estimators that support multioutput regression may be faster.
- E.g., prediction of both wind speed and wind direction, in degrees, using data obtained at a certain location.
- Each sample would be data obtained at one location and both wind speed and direction would be output for each sample.



Comparing Learning Systems

- How can we determine if one learning system provides better performance than another
- Example:

| | <u>Accuracies on test sets</u> | | | | |
|------------|--------------------------------|----|----|-----|----|
| System 1: | 80% | 50 | 75 | ... | 99 |
| System 2: | 79 | 49 | 74 | ... | 98 |
| δ : | +1 | +1 | +1 | ... | +1 |

- Notice that System 1 is always better than System 2

Hypothesis Testing

- Hypothesis testing (or test of significance) is a statistical method used to determine if there is enough evidence in a sample data to draw conclusions about a population.
- Given a claim, identify the
 - **The null hypothesis (H_0)** = the value of a population parameter (e.g., mean, or SD) is equal to some claimed value.
 - **The alternative hypothesis (H_a)** the parameter differs from H_0
- Standard test for testing the difference between population:
paired t-test and sign test

Comparing Systems Using a Paired t-Test

- consider δ 's as observed values of a set of i.i.d. random variables
- H_0 : the 2 learning systems have the same accuracy
- H_a : one of the systems is more accurate than the other
- Use paired t-test to compute p-value or significance level that mean of δ 's would arise from null hypothesis
- If p-value is sufficiently small (typically < 0.05) then reject the null hypothesis

Comparing Systems Using a Paired t-Test


- Calculate the t statistic
- Determine the corresponding p-value, by looking up t in a table of values for the Student's t-distribution with n-1 degrees of freedom

$$t = \frac{\bar{\delta}}{\sqrt{\frac{1}{n(n-1)} \sum_{i=1}^n (\delta_i - \bar{\delta})^2}}$$

$$\bar{\delta} = \frac{1}{n} \sum_{i=1}^n \delta_i$$

Comparing Systems Using a Paired t-Test


- The null distribution of our t statistic looks like this
- The p-value indicates how far out in a tail our t statistic is
- If the p-value is sufficiently small, we reject the null hypothesis, since it is unlikely we'd get such a value by chance



for a two-tailed test, the p -value represents the probability mass in these two regions

Why Do We Use a Two-Tailed Test?

- A two-tailed test asks the question: is the accuracy of the two systems different?
- A one-tailed test asks the question: is system A better than system B?
- A priori, we don't know which learning system will be more accurate (if there is a difference) – we want to allow that either one might be



One-tailed Test Vs Two-tailed Test

Sign Test

- Test the null hypothesis that the **median of a distribution is equal to some value.**
- It is a non-parametric or “distribution-free” test = Don’t assume that data comes from a particular distribution, like the normal distribution.
- Count “wins” for Algorithm A and B over the N test examples on which they disagree
- Let M be the larger of these counts
- What is probability under $b(N, 0.5)$ that either A or B would win at least M times? → Find p-value from binomial table.

The top corners of the slide feature decorative circuit-like patterns. These consist of thin, light blue lines that branch out and connect to small circular nodes, resembling a stylized electronic circuit or network diagram.

2.

Case Studies

Case Study : Kaggle

- Look at competition (maybe completed one) and notice:
 - Evaluation (Hold-out)
 - Metrics
 - How people get the solution (<https://medium.com/kaggle-blog>)
 - Interviews with Machine Learning Heroes (Kagglers + Practitioners + Researchers) (<https://www.kaggle.com/discussions/general/76241>)

Kaggle Winner's Blog
Interviews with top Kagglers by the Kaggle Team



OTTO – Multi-Objective Recommender System

- Build a multi-objective recommender system based on real-world e-commerce sessions
→ predict e-commerce clicks, cart additions, and orders.
- <https://www.kaggle.com/competitions/otto-recommender-system>
- Evaluation: Evaluated on the first 20 predictions for each action type, and the three recall values are weight-averaged:

$$score = 0.10 \cdot R_{clicks} + 0.30 \cdot R_{carts} + 0.60 \cdot R_{orders}$$

$$R_{type} = \frac{\sum_i^N |\{\text{predicted aids}\}_{i,type} \cap \{\text{ground truth aids}\}_{i,type}|}{\sum_i^N \min(20, |\{\text{ground truth aids}\}_{i,type}|)}$$

American Express - Default Prediction

- Predict if a customer will default in the future
- <https://www.kaggle.com/competitions/amex-default-prediction>
- The evaluation metric (M) = the mean of two measures of rank ordering: Normalized Gini Coefficient (G) and default rate captured at 4% (D):

$$M = 0.5 \cdot (G + D)$$

Kaggle - LLM Science Exam

- Answer difficult science-based questions written by a Large Language Model (LLM).
- <https://www.kaggle.com/competitions/kaggle-llm-science-exam>
- Submissions are evaluated according to the Mean Average Precision @ 3 (MAP@3):

$$MAP@3 = \frac{1}{U} \sum_{u=1}^U \sum_{k=1}^{\min(n,3)} P(k) \times rel(k)$$

- This competition uses a hidden test.

Store Sales - Time Series Forecasting

- Use machine learning to predict grocery sales
- <https://www.kaggle.com/competitions/store-sales-time-series-forecasting>
- The evaluation metric for this competition is **Root Mean Squared Logarithmic Error**.

$$\sqrt{\frac{1}{n} \sum_{i=1}^n (\log(1 + \hat{y}_i) - \log(1 + y_i))^2}$$

- How to split train-test data for time-series

NFL Big Data Bowl

- How many yards will an NFL player gain after receiving a handoff?
- <https://www.kaggle.com/competitions/nfl-big-data-bowl-2020>
- Submissions will be evaluated on the Continuous Ranked Probability Score (CRPS)
→ Predict the probability that the team gains \leq that many yards on the play

$$C = \frac{1}{199N} \sum_{m=1}^N \sum_{n=-99}^{99} (P(y \leq n) - H(n - Y_m))^2$$

- **Solution motivation:** <https://www.kaggle.com/c/nfl-big-data-bowl-2020/discussion/119400>

Exercise:

Binary Classification with a Software Defects Dataset

- <https://www.kaggle.com/competitions/playground-series-s3e23>
- Predict defects in C programs given various attributes about the code.
- Feature description:
<https://www.kaggle.com/datasets/semustafacevik/software-defect-prediction>

% 7. Attribute Information:

```
%  
%  
% 1. loc : numeric % McCabe's line count of code  
% 2. v(g) : numeric % McCabe "cyclomatic complexity"  
% 3. ev(g) : numeric % McCabe "essential complexity"  
% 4. iv(g) : numeric % McCabe "design complexity"  
% 5. n : numeric % Halstead total operators + operands  
% 6. v : numeric % Halstead "volume"  
% 7. l : numeric % Halstead "program length"  
% 8. d : numeric % Halstead "difficulty"  
% 9. i : numeric % Halstead "intelligence"  
% 10. e : numeric % Halstead "effort"  
% 11. b : numeric % Halstead  
% 12. t : numeric % Halstead's time estimator  
% 13. l0Code : numeric % Halstead's line count  
% 14. l0Comment : numeric % Halstead's count of lines of comments  
% 15. l0Blank : numeric % Halstead's count of blank lines  
% 16. l0CodeAndComment: numeric  
% 17. uniq_Op : numeric % unique operators  
% 18. uniq_Opnd : numeric % unique operands  
% 19. total_Op : numeric % total operators  
% 20. total_Opnd : numeric % total operands  
% 21. branchCount : numeric % of the flow graph  
% 22. defects : {false,true} % module has/has not one or more  
% % reported defects
```